

```
In [1]: import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import seaborn as sns
from sklearn.linear_model import LogisticRegression
from sklearn.model_selection import train_test_split
```

```
In [2]: df=pd.read_csv(r"C4_framingham.csv")
df
```

```
Out[2]:
```

	male	age	education	currentSmoker	cigsPerDay	BPMeds	prevalentStroke	prevalentHyp
0	1	39	4.0	0	0.0	0.0	0	0
1	0	46	2.0	0	0.0	0.0	0	0
2	1	48	1.0	1	20.0	0.0	0	0
3	0	61	3.0	1	30.0	0.0	0	1
4	0	46	3.0	1	23.0	0.0	0	0
...
4233	1	50	1.0	1	1.0	0.0	0	1
4234	1	51	3.0	1	43.0	0.0	0	0
4235	0	48	2.0	1	20.0	NaN	0	0
4236	0	44	1.0	1	15.0	0.0	0	0
4237	0	52	2.0	0	0.0	0.0	0	0

4238 rows × 16 columns



In [3]: df.info()

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 4238 entries, 0 to 4237
Data columns (total 16 columns):
#   Column                Non-Null Count  Dtype
---  -
0   male                  4238 non-null   int64
1   age                   4238 non-null   int64
2   education             4133 non-null   float64
3   currentSmoker         4238 non-null   int64
4   cigsPerDay            4209 non-null   float64
5   BPMeds                4185 non-null   float64
6   prevalentStroke       4238 non-null   int64
7   prevalentHyp          4238 non-null   int64
8   diabetes              4238 non-null   int64
9   totChol               4188 non-null   float64
10  sysBP                 4238 non-null   float64
11  diaBP                 4238 non-null   float64
12  BMI                   4219 non-null   float64
13  heartRate             4237 non-null   float64
14  glucose               3850 non-null   float64
15  TenYearCHD            4238 non-null   int64
dtypes: float64(9), int64(7)
memory usage: 529.9 KB
```

In [4]: df=df.dropna()

In [5]: df.describe()

Out[5]:

	male	age	education	currentSmoker	cigsPerDay	BPMeds	prevalen
count	3656.000000	3656.000000	3656.000000	3656.000000	3656.000000	3656.000000	3656
mean	0.443654	49.557440	1.979759	0.489059	9.022155	0.030361	(
std	0.496883	8.561133	1.022657	0.499949	11.918869	0.171602	(
min	0.000000	32.000000	1.000000	0.000000	0.000000	0.000000	(
25%	0.000000	42.000000	1.000000	0.000000	0.000000	0.000000	(
50%	0.000000	49.000000	2.000000	0.000000	0.000000	0.000000	(
75%	1.000000	56.000000	3.000000	1.000000	20.000000	0.000000	(
max	1.000000	70.000000	4.000000	1.000000	70.000000	1.000000	✓

In [7]: df.columns

Out[7]: Index(['male', 'age', 'education', 'currentSmoker', 'cigsPerDay', 'BPMeds', 'prevalentStroke', 'prevalentHyp', 'diabetes', 'totChol', 'sysBP', 'diaBP', 'BMI', 'heartRate', 'glucose', 'TenYearCHD'], dtype='object')

```
In [8]: df1=df[['male', 'age', 'education', 'currentSmoker', 'cigsPerDay', 'BPMeds',  
              'prevalentStroke', 'prevalentHyp', 'diabetes', 'totChol', 'sysBP',  
              'diaBP', 'BMI', 'heartRate', 'glucose', 'TenYearCHD']]
```

```
In [20]: x=df1[['age', 'education', 'currentSmoker', 'cigsPerDay', 'BPMeds',  
              'prevalentStroke', 'prevalentHyp', 'diabetes', 'totChol', 'sysBP',  
              'diaBP', 'BMI', 'heartRate', 'glucose', 'TenYearCHD']]  
y=df1['male']  
x_train,x_test,y_train,y_test=train_test_split(x,y,test_size=0.3)  
lr=LogisticRegression()  
lr.fit(x_train,y_train)
```

C:\ProgramData\Anaconda3\lib\site-packages\sklearn\linear_model_logistic.py:
763: ConvergenceWarning: lbfgs failed to converge (status=1):
STOP: TOTAL NO. of ITERATIONS REACHED LIMIT.

Increase the number of iterations (max_iter) or scale the data as shown in:
<https://scikit-learn.org/stable/modules/preprocessing.html> (<https://scikit-learn.org/stable/modules/preprocessing.html>)
Please also refer to the documentation for alternative solver options:
https://scikit-learn.org/stable/modules/linear_model.html#logistic-regression (https://scikit-learn.org/stable/modules/linear_model.html#logistic-regression)
n_iter_i = _check_optimize_result(

Out[20]: LogisticRegression()

```
In [21]: lr.predict(x_test)
```

Out[21]: array([0, 0, 0, ..., 0, 0, 0], dtype=int64)

```
In [22]: lr.score(x_test,y_test)
```

Out[22]: 0.6918869644484958

```
In [23]: from sklearn.preprocessing import StandardScaler  
fs=StandardScaler().fit_transform(x)  
logr=LogisticRegression()  
logr.fit(fs,y)
```

Out[23]: LogisticRegression()

```
In [24]: o=[[1,2,3,4,5,6,7,8,9,10,11,12,13,14,15]]  
prediction=logr.predict(o)  
print(prediction)
```

[1]

```
In [25]: logr.classes_
```

Out[25]: array([0, 1], dtype=int64)

```
In [26]: logr.predict_proba(o)[0][0]
```

```
Out[26]: 0.023655703523128735
```

```
In [27]: logr.predict_proba(o)[0][1]
```

```
Out[27]: 0.9763442964768713
```

```
In [ ]:
```